

```
In [46]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import datasets
```

```
In [47]: # Loading the California Housing Dataset
housing = datasets.fetch_california_housing()
```

```
In [48]: cali_housing = housing.data
```

```
In [49]: features = housing.feature_names
```

```
In [50]: features
```

```
Out[50]: ['MedInc',
          'HouseAge',
          'AveRooms',
          'AveBedrms',
          'Population',
          'AveOccup',
          'Latitude',
          'Longitude']
```

```
In [51]: cali_housing.shape
```

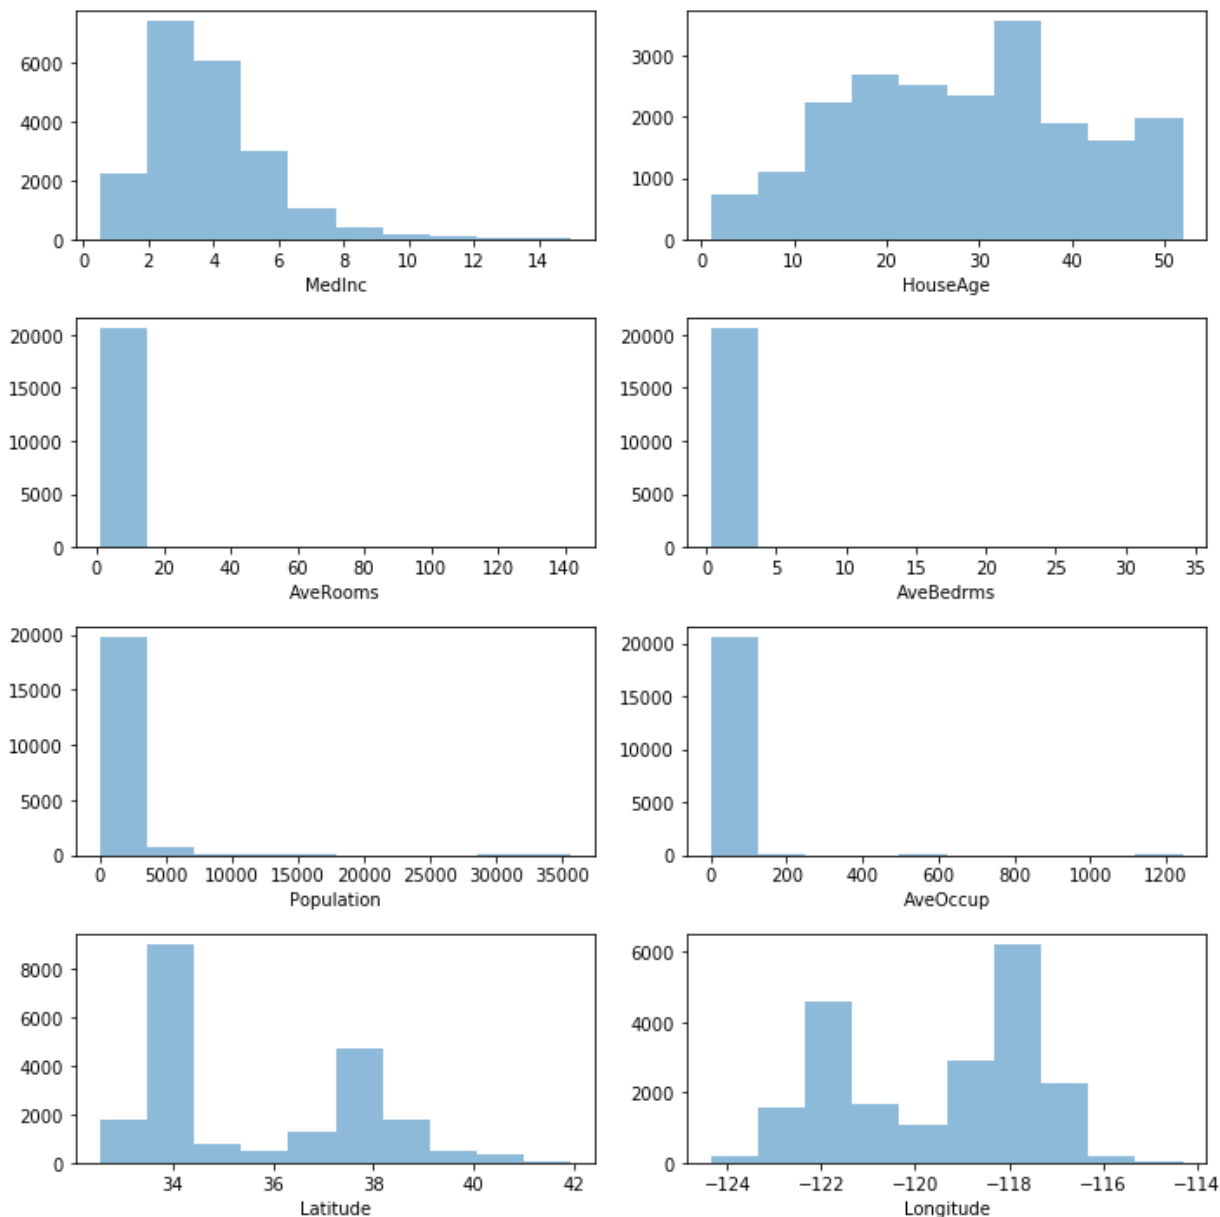
```
Out[51]: (20640, 8)
```

```
In [52]: median_house_value = housing.target
```

## 1.1


In [53]: *# Visualizing univariate distribution of each feature*

```
count = 0
fig = plt.figure(figsize=(10,10))
for i in range(1,9):
    fig.add_subplot(4,2,i)
    plt.hist(cali_housing[:,count], alpha=0.5)
    plt.xlabel(features[count])
    #plt.ylabel("Median House Value")
    plt.tight_layout()
    count += 1
plt.show()
```



```
In [56]: # Visualizing univariate distribution of target

plt.hist(median_house_value, alpha=0.5)
plt.xlabel("Median House Value")
plt.show()
```



### Observations:

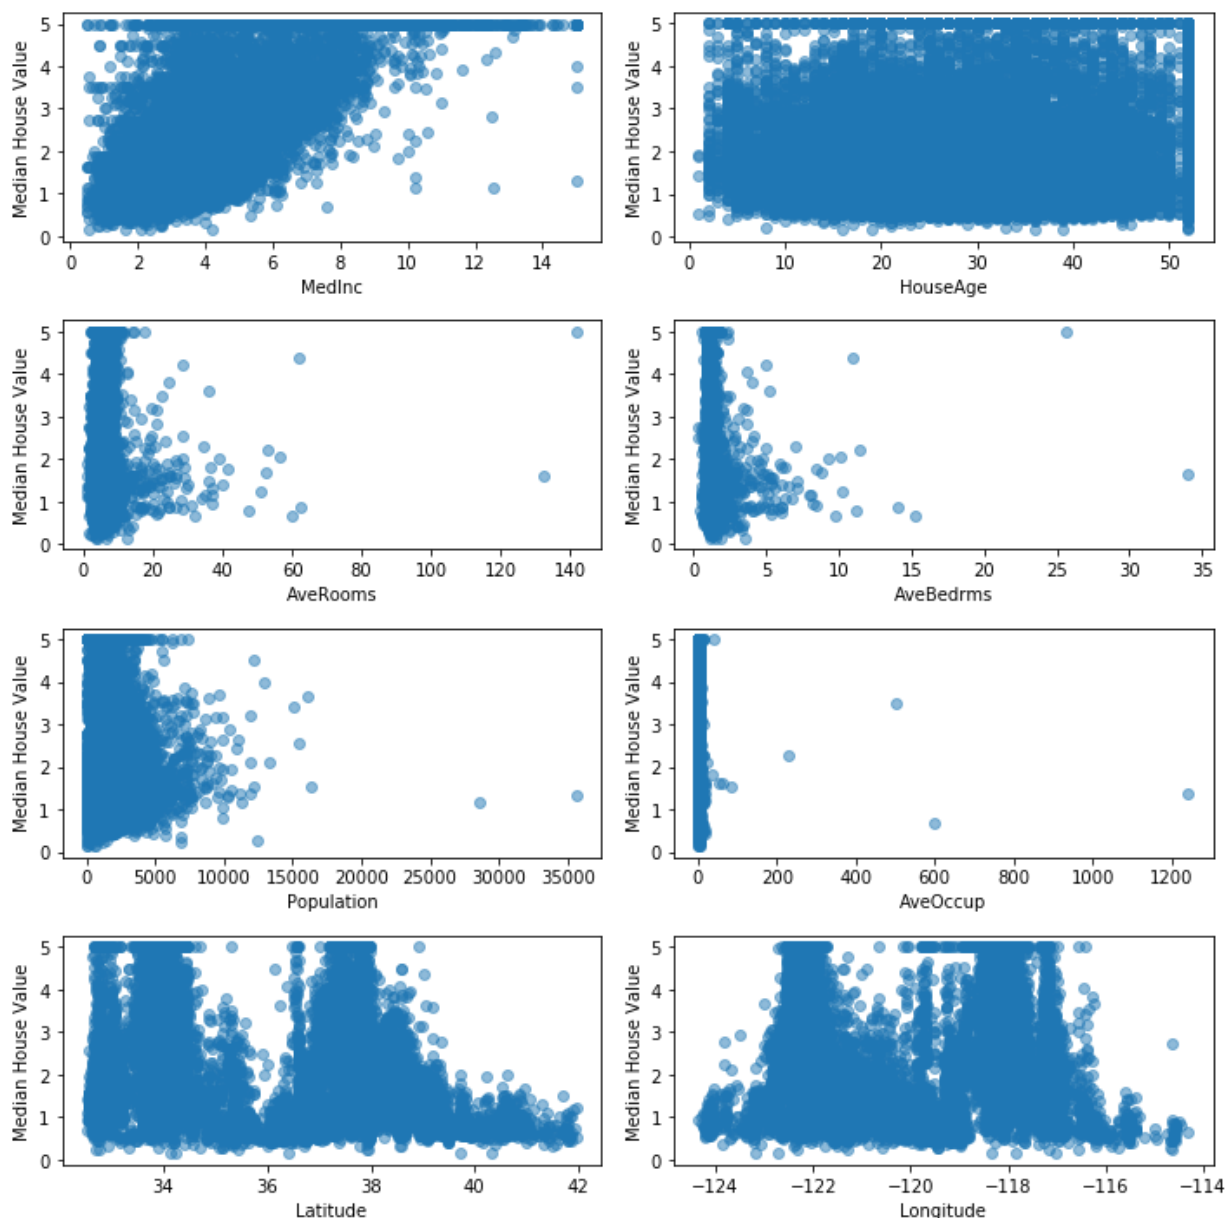
- 1) There are outliers in Average Occupancy of houses
- 2) Latitude and Longitude are bimodal which suggests that there are a couple of densely populated areas in California

## 1.2

```
In [57]: # Visualizing the dependency of target on each feature

count = 0
fig = plt.figure(figsize=(10,10))
for i in range(1,9):
    fig.add_subplot(4,2,i)
    plt.scatter(cali_housing[:,count], median_house_value, alpha=0.5)
    plt.xlabel(features[count])
    plt.ylabel("Median House Value")
    plt.tight_layout()
    count += 1
```

```
In [58]: plt.show()
```



### 1.3

```
In [59]: from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
```

```
In [60]: X_train, X_test, y_train, y_test = train_test_split(
    cali_housing, median_house_value, random_state=0)
```

```
In [61]: models = [LinearRegression(), Ridge(), Lasso(), ElasticNet()]
```

```
In [62]: models_names = ['LinearRegression', 'Ridge', 'Lasso', 'ElasticNet']
```

```
In [64]: for i in range(len(models)):
          score_train = np.mean(cross_val_score(models[i], X_train, y_train, cv=10))
          print("Mean training score of " + models_names[i] + " : " + str(score_train))
```

```
Mean training score of LinearRegression : 0.606136752898
Mean training score of Ridge : 0.606151468466
Mean training score of Lasso : 0.291549120926
Mean training score of ElasticNet : 0.429496808434
```

```
In [65]: # Scaling the data with Standard Scaler
```

```
scaler = StandardScaler()
scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [66]: for i in range(len(models)):
          score_train_scaled = np.mean(cross_val_score(models[i], X_train_scaled, y_train_scaled, cv=10))
          print("Mean training score of " + models_names[i] + " : " + str(score_train_scaled))
```

```
Mean training score of LinearRegression : 0.606136752898
Mean training score of Ridge : 0.606143117729
Mean training score of Lasso : -0.000601150113709
Mean training score of ElasticNet : 0.209910128135
```

Observation: scaling the data with StandardScaler doesn't help

## 1.4

```
In [67]: param = [{'alpha': np.logspace(-3, 3, 13)}, {'alpha': np.logspace(-3, 0, 13)}]
          'l1_ratio': [0.01, .1, .5, .9, .98, 1]]
          print(param)
```

```
[{'alpha': array([ 1.00000000e-03,  3.16227766e-03,  1.00000000e-02,
                  3.16227766e-02,  1.00000000e-01,  3.16227766e-01,
                  1.00000000e+00,  3.16227766e+00,  1.00000000e+01,
                  3.16227766e+01,  1.00000000e+02,  3.16227766e+02,
                  1.00000000e+03])), {'alpha': array([ 0.001, 0.00177828,
                  0.00316228, 0.00562341, 0.01, 0.01778279,
                  0.03162278, 0.05623413, 0.1, 0.17782794,
                  0.31622777, 0.56234133, 1.]), 'l1_ratio': [0.01, 0.1,
                  0.5, 0.9, 0.98, 1]}, {'alpha': array([ 0.0001, 0.00021544, 0.00046416,
                  0.001, 0.00215443, 0.00464159, 0.01, 0.02154435, 0.04641589, 0.1,
                  0.21544347, 0.46415888, 1.]), 'l1_ratio': [0.01, 0.1,
                  0.5, 0.9, 0.98, 1]}]
```

```
In [113]: # Tuning the parameters of the models using GridSearchCV
list_of_result_dataframes = list()
for i in range(len(models)-1):
    grid = GridSearchCV(models[i+1], param[i], cv=10)
    grid.fit(X_train, y_train)
    list_of_result_dataframes.append(pd.DataFrame(grid.cv_results_))
```

```
In [114]: # Results for Grid Search on Ridge
list_of_result_dataframes[0]
```

```
Out[114]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_alpha	param_alpha__alpha_1
0	0.001910	0.000502	0.606137	0.611152	0.001	{'alpha': 0.001}
1	0.001699	0.000280	0.606137	0.611152	0.00316228	{'alpha': 0.0031622776601683793}
2	0.002480	0.000433	0.606137	0.611152	0.01	{'alpha': 0.01}
3	0.002957	0.000491	0.606137	0.611152	0.0316228	{'alpha': 0.031622776601683793}
4	0.001602	0.000277	0.606138	0.611152	0.1	{'alpha': 0.1}
5	0.001261	0.000210	0.606141	0.611152	0.316228	{'alpha': 0.31622776601683793}
6	0.001223	0.000207	0.606151	0.611151	1	{'alpha': 1}
7	0.001208	0.000204	0.606183	0.611151	3.16228	{'alpha': 3.1622776601683793}
8	0.001204	0.000203	0.606274	0.611146	10	{'alpha': 10}
9	0.001598	0.000283	0.606504	0.611104	31.6228	{'alpha': 31.622776601683793}
10	0.001210	0.000217	0.606807	0.610777	100	{'alpha': 100}
11	0.001296	0.000227	0.605999	0.609052	316.228	{'alpha': 316.22776601683793}
12	0.001229	0.000222	0.600725	0.603141	1000	{'alpha': 1000}

13 rows × 31 columns

```
In [115]: # Results for Grid Search on Lasso
list_of_result_dataframes[1]
```

```
Out[115]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_alpha	param_epsilon
0	0.009053	0.000232	0.606374	0.611114	0.001	{'alpha': 0.001, 'epsilon': 0.001}
1	0.008635	0.000231	0.606498	0.611033	0.00177828	{'alpha': 0.0017782794, 'epsilon': 0.0017782794}
2	0.009546	0.000295	0.606589	0.610777	0.00316228	{'alpha': 0.0031622776, 'epsilon': 0.0031622776}
3	0.009629	0.000356	0.606337	0.609967	0.00562341	{'alpha': 0.0056234133, 'epsilon': 0.0056234133}
4	0.008558	0.000255	0.604583	0.607406	0.01	{'alpha': 0.01, 'epsilon': 0.01}
5	0.006588	0.000794	0.600121	0.602024	0.0177828	{'alpha': 0.017782794, 'epsilon': 0.017782794}
6	0.005416	0.000228	0.594291	0.595973	0.0316228	{'alpha': 0.031622776, 'epsilon': 0.031622776}
7	0.003864	0.000238	0.583396	0.584626	0.0562341	{'alpha': 0.056234133, 'epsilon': 0.056234133}
8	0.003424	0.000226	0.549299	0.550440	0.1	{'alpha': 0.1, 'epsilon': 0.1}
9	0.001455	0.000220	0.513368	0.514470	0.177828	{'alpha': 0.17782794, 'epsilon': 0.17782794}
10	0.001450	0.000262	0.494713	0.495351	0.316228	{'alpha': 0.31622776, 'epsilon': 0.31622776}
11	0.001385	0.000223	0.445874	0.446528	0.562341	{'alpha': 0.56234133, 'epsilon': 0.56234133}
12	0.001390	0.000239	0.291549	0.292297	1	{'alpha': 1, 'epsilon': 1}

13 rows × 31 columns

```
In [116]: # Results for Grid Search on Elastic Net  
list_of_result_dataframes[2]
```

<b>41</b>	0.007717	0.000231	0.604583	0.607406	0.01	1
<b>42</b>	0.008841	0.000252	0.606061	0.609142	0.0215443	0.01
<b>43</b>	0.007478	0.000231	0.605683	0.608647	0.0215443	0.1
<b>44</b>	0.006884	0.000235	0.602653	0.605055	0.0215443	0.5
<b>45</b>	0.005949	0.000234	0.599351	0.601212	0.0215443	0.9
<b>46</b>	0.005848	0.000234	0.598845	0.600683	0.0215443	0.98
---	-----	-----	-----	-----	-----	-



```


In [134]: save_images = ['Ridge.png', 'Lasso.png']
for i in range(0,2):
    plt.figure()
    ax1 = plt.gca()
    line1, = ax1.plot(list_of_result_dataframes[i]['param_alpha'], list_of_result_dataframes[i]['mean_train_score'])
    line2, = ax1.plot(list_of_result_dataframes[i]['param_alpha'], list_of_result_dataframes[i]['mean_test_score'])

    plt.fill_between(list_of_result_dataframes[i].param_alpha.astype(np.float64),
                     list_of_result_dataframes[i]['mean_train_score'] + list_of_result_dataframes[i]['std_train_score'],
                     list_of_result_dataframes[i]['mean_train_score'] - list_of_result_dataframes[i]['std_train_score'],
                     color='lightblue')
    plt.fill_between(list_of_result_dataframes[i].param_alpha.astype(np.float64),
                     list_of_result_dataframes[i]['mean_test_score'] + list_of_result_dataframes[i]['std_test_score'],
                     list_of_result_dataframes[i]['mean_test_score'] - list_of_result_dataframes[i]['std_test_score'],
                     color='lightred')

    plt.legend([line1, line2], ["mean_train_score", "mean_test_score"], loc='best')
    ax1.set_ylabel("score")
    ax1.set_xlabel("alpha")
    ax1.set_xscale("log")
    if i==0:
        plt.title("score vs alpha for Ridge Regression")
    else:
        plt.title("score vs alpha for Lasso Regression")

    plt.savefig(save_images[i], bbox_inches = 'tight')
    plt.show()

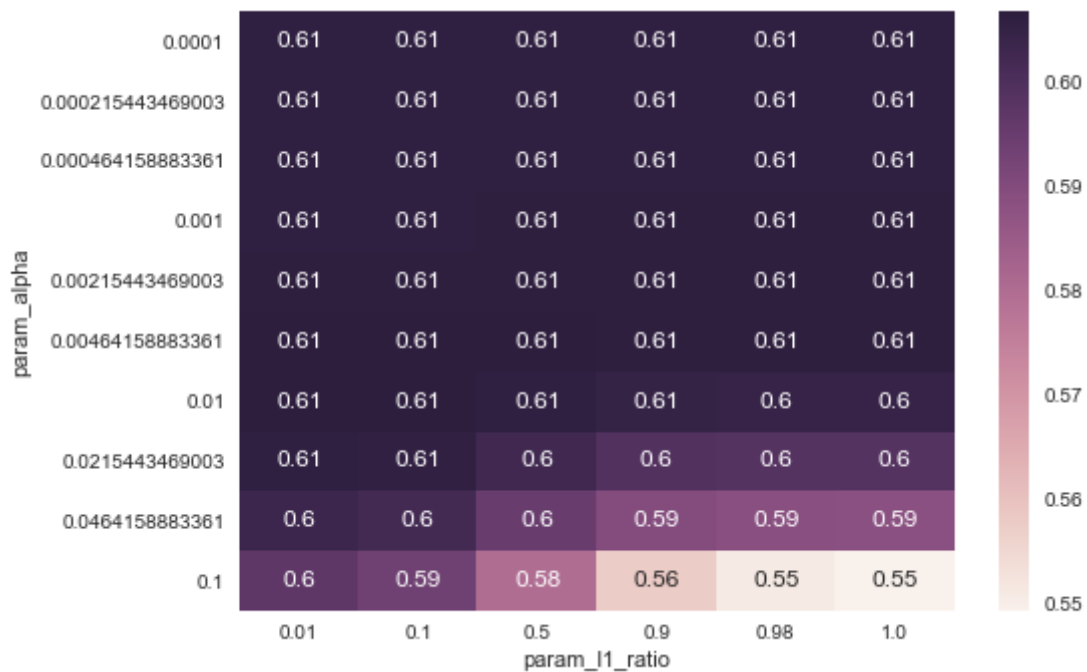
```





```
In [135]: res = pd.pivot_table(list_of_result_dataframes[2],
    values='mean_test_score', index='param_alpha', columns='param_l1_ratio')

import seaborn as sns; sns.set()
sns.heatmap(res, annot = True)
plt.show()
```



In [119]: res

```

Out[119]:
  param_l1_ratio    0.01    0.1    0.5    0.9    0.98    1.0
  param_alpha
0.000100  0.606157  0.606158  0.606161  0.606164  0.606164  0.606164
0.000215  0.606180  0.606182  0.606188  0.606194  0.606195  0.606195
0.000464  0.606228  0.606231  0.606243  0.606255  0.606257  0.606258
0.001000  0.606323  0.606328  0.606350  0.606369  0.606373  0.606374
0.002154  0.606490  0.606498  0.606524  0.606538  0.606539  0.606539
0.004642  0.606716  0.606715  0.606671  0.606547  0.606510  0.606501
0.010000  0.606787  0.606716  0.606155  0.605007  0.604673  0.604583
0.021544  0.606061  0.605683  0.602653  0.599351  0.598845  0.598710
0.046416  0.603436  0.602144  0.595231  0.589820  0.588683  0.588383
0.100000  0.597221  0.593782  0.580172  0.557644  0.551079  0.549299

```

Observations:

There is significant improvement in results for Lasso and Elastic Net whereas for Ridge regression the results are almost the same.


## 1.5



```

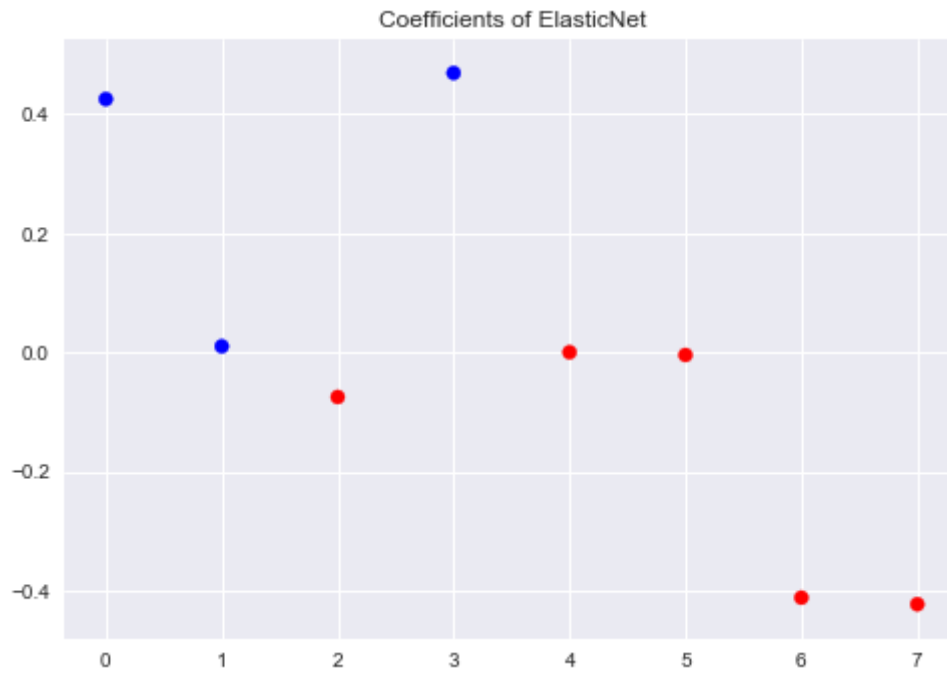
In [133]: save_images_names = ['LinearRegression_Scatter', 'Ridge_Scatter', 'Lasso_Scatter']
for i in range(len(models)):
    if i==0:
        grid = LinearRegression().fit(X_train, y_train, sample_weight=None)
        plt.scatter(range(8), grid.coef_,
                    c=np.sign(grid.coef_), cmap="bwr_r")
        plt.title("Coefficients of " + str(models_names[i]))
        plt.savefig(save_images_names[i], bbox_inches = 'tight')
        plt.show()
    else:
        grid = GridSearchCV(models[i], param[i-1], cv=10)
        grid.fit(X_train, y_train)

        plt.scatter(range(8), grid.best_estimator_.coef_,
                    c=np.sign(grid.best_estimator_.coef_), cmap="bwr_r")
        plt.title("Coefficients of " + str(models_names[i]))
        plt.savefig(save_images_names[i], bbox_inches = 'tight')
        plt.show()

```







Observation: All the models agree on which features are important

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import datasets
```

```
In [2]: from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
```

## 2.1

```
In [3]: forest = datasets.fetch_covtype()
```

```
In [4]: forest_data = forest.data
#forest_features = forest.feature_names
cover_type = forest.target
```

```
In [5]: forest_data.shape
```

```
Out[5]: (581012, 54)
```

```
In [6]: forest_features = ['elevation', 'aspect', 'slope', 'hdist_to_hydrology', 'vdist_to_hydrology',
                           'hill_shade_9', 'hill_shade_12', 'hill_shade_3', 'hdist_to_fire']
```

```
In [7]: for i in range(1,5):
        forest_features.append('wilderness' + str(i))
```

```
In [8]: forest_features
```

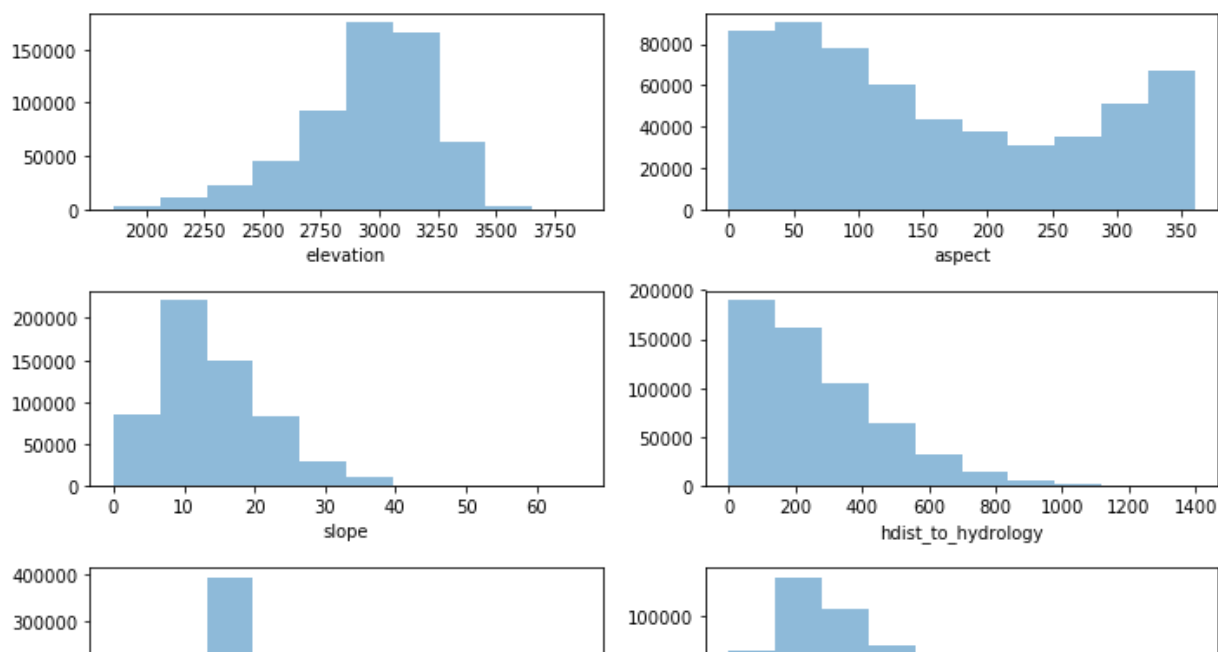
```
Out[8]: ['elevation',
         'aspect',
         'slope',
         'hdist_to_hydrology',
         'vdist_to_hydrology',
         'hdist_to_road',
         'hill_shade_9',
         'hill_shade_12',
         'hill_shade_3',
         'hdist_to_fire',
         'wilderness1',
         'wilderness2',
         'wilderness3',
         'wilderness4']
```

```
In [9]: for i in range(1,41):
        forest_features.append('soil_type' + str(i))
```

```
In [10]: len(forest_features)
```

```
Out[10]: 54
```

```
In [11]: # Visualizing univariate distribution of each feature
count = 0
fig = plt.figure(figsize=(10,60))
for i in range(1,55):
    fig.add_subplot(27,2,i)
    plt.hist(forest_data[:,count], alpha=0.5)
    plt.xlabel(forest_features[count])
    #plt.ylabel("Median House Value")
    plt.tight_layout()
    count += 1
plt.show()
```




```
In [12]: count
```

```
Out[12]: 54
```



```
In [13]: # Visualizing univariate distribution of target
plt.hist(cover_type, alpha=0.5, bins=8)
plt.xlabel('Cover_type')
plt.show()
```



## 2.2

```
In [14]: X_train, X_test, y_train, y_test = train_test_split(
        forest_data, cover_type, random_state=0)
```

```
In [15]: from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
        from sklearn.svm import LinearSVC
        from sklearn.neighbors.nearest_centroid import NearestCentroid
```

```
In [44]: #logreg = LogisticRegressionCV(multi_class='multinomial', dual=False).fit(X_
```

```
In [45]: # Accuracy over train data
        #print(logreg.score(X_train,y_train))
```

```
In [16]: cross_validation_scores_logreg = cross_val_score(LogisticRegressionCV(multi_
```

```
In [22]: mean_cross_validation_score_logreg = np.mean(cross_validation_scores_logreg)
        print("mean cross validation score of logreg is " + str(mean_cross_validation_score_logreg))

mean cross validation score of logreg is 0.672098597965
```

```
In [23]: cross_validation_scores_linearsvm = cross_val_score(LinearSVC(dual=False, tol=
```

```
In [25]: mean_cross_validation_score_linearsvm = np.mean(cross_validation_scores_linearsvm)
        print("mean cross validation score of linearsvm is " + str(mean_cross_validation_score_linearsvm))

mean cross validation score of linearsvm is 0.680220061782
```

```
In [27]: cross_validation_scores_nearest_centroid = cross_val_score(NearestCentroid(),
```

```
In [28]: mean_cross_validation_score_nearest_centroid = np.mean(cross_validation_scores_nearest_centroid)
print("mean cross validation score of nearest centroid is " + str(mean_cross_validation_score_nearest_centroid))

mean cross validation score of nearest centroid is 0.193586356619
```

```
In [29]: # Scaling the data with StandardScaler
scaler = StandardScaler()
scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [30]: cross_validation_scores_logreg = cross_val_score(LogisticRegressionCV(multi_class='ovr'), X_train_scaled, y_train_scaled, cv=5)
mean_cross_validation_score_logreg = np.mean(cross_validation_scores_logreg)
print("mean cross validation score of logreg is " + str(mean_cross_validation_score_logreg))

mean cross validation score of logreg is 0.72434305351
```

```
In [31]: cross_validation_scores_linearsvm = cross_val_score(LinearSVC(dual=False, tol=1e-5), X_train_scaled, y_train_scaled, cv=5)
mean_cross_validation_score_linearsvm = np.mean(cross_validation_scores_linearsvm)
print("mean cross validation score of linearsvm is " + str(mean_cross_validation_score_linearsvm))

mean cross validation score of linearsvm is 0.712719657497
```

```
In [32]: cross_validation_scores_nearest_centroid = cross_val_score(NearestCentroid(), X_train_scaled, y_train_scaled, cv=5)
mean_cross_validation_score_nearest_centroid = np.mean(cross_validation_scores_nearest_centroid)
print("mean cross validation score of nearest centroid is " + str(mean_cross_validation_score_nearest_centroid))

mean cross validation score of nearest centroid is 0.549861256828
```

Observations:

- 1) Logistic Regression and Linear Support Vector Machines are giving significantly higher cross validation scores than Nearest Centroid. (Best Score: linear support vector machine: 0.68)
- 2) Scaling the data using StandardScaler gives better cross validation score than without scaling the data in case of all three models and even in this case, logistic regression and linear support vector machines are performing better than nearest centroid. (Best Score: logistic regression: 0.724)

As the models are performing better when X\_train is scaled, we use the standard scaled X\_data from now.

## 2.3

```
In [34]: params = [{'C': [0.01, 1, 10, 100, 1000]}]
grid_logreg = GridSearchCV(LogisticRegression(multi_class='multinomial', dual=True),
                             params, cv=5)
grid_logreg.fit(X_train_scaled, y_train)
pd.DataFrame(grid_logreg.cv_results_)
```

```
Out[34]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	params	rank_test_score
0	14.930232	0.029203	0.722741	0.722941	0.01	{'C': 0.01}	
1	14.429111	0.027336	0.723873	0.724003	1	{'C': 1}	
2	14.432691	0.029546	0.724065	0.724046	10	{'C': 10}	
3	14.886415	0.027300	0.724081	0.723992	100	{'C': 100}	
4	15.148282	0.027407	0.723958	0.723974	1000	{'C': 1000}	

```
In [35]: params = [{'C': [0.01, 1, 10, 100, 1000]}]
grid_linearsvc = GridSearchCV(LinearSVC(dual=False, tol=0.001), params, cv=5)
grid_linearsvc.fit(X_train_scaled, y_train)
pd.DataFrame(grid_linearsvc.cv_results_)
```

```
Out[35]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	params	rank_test_score
0	37.866422	0.027945	0.712424	0.712354	0.01	{'C': 0.01}	
1	40.058371	0.028552	0.712720	0.712685	1	{'C': 1}	
2	39.194059	0.028445	0.712738	0.712701	10	{'C': 10}	
3	39.344713	0.028619	0.712736	0.712707	100	{'C': 100}	
4	38.208435	0.027860	0.712736	0.712699	1000	{'C': 1000}	

```
In [50]: params = [{'shrink_threshold': [0,0.5,1,10, 15, 20, 50, 100]}]
grid_nearest_centroid = GridSearchCV(NearestCentroid(), params, cv=3)
grid_nearest_centroid.fit(X_train_scaled,y_train)
pd.DataFrame(grid_nearest_centroid.cv_results_)
```

```
Out[50]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_shrink_threshold
<b>0</b>	0.187555	0.036439	0.549861	0.549946	0 {'st
<b>1</b>	0.344647	0.036744	0.548528	0.548660	0.5 {'st
<b>2</b>	0.341735	0.035756	0.547204	0.547211	1 {'st
<b>3</b>	0.342701	0.035268	0.542938	0.543083	10 {'st
<b>4</b>	0.339407	0.034746	0.547222	0.547348	15 {'st
<b>5</b>	0.338873	0.034924	0.552881	0.552903	20 {'st
<b>6</b>	0.336124	0.034707	0.633974	0.633935	50 {'st
<b>7</b>	0.335678	0.035027	0.514564	0.514576	100 {'st

Observation: There is a significant improvement in case of Nearest Centroid whereas the others are pretty much the same.


```

In [42]: plt.figure()
ax1 = plt.gca()
logreg_dataframe = pd.DataFrame(grid_logreg.cv_results_)
line1, = ax1.plot(logreg_dataframe['param_C'], logreg_dataframe['mean_train_score'])
line2, = ax1.plot(logreg_dataframe['param_C'], logreg_dataframe['mean_test_score'])

plt.fill_between(logreg_dataframe.param_C.astype(np.float),
                 logreg_dataframe['mean_train_score'] + logreg_dataframe['std_train_score'],
                 logreg_dataframe['mean_train_score'] - logreg_dataframe['std_train_score'],
                 color='gray')
plt.fill_between(logreg_dataframe.param_C.astype(np.float),
                 logreg_dataframe['mean_test_score'] + logreg_dataframe['std_test_score'],
                 logreg_dataframe['mean_test_score'] - logreg_dataframe['std_test_score'],
                 color='orange')

plt.legend([line1, line2], ["mean_train_score", "mean_test_score"], loc="upper right")
ax1.set_ylabel("score")
ax1.set_xlabel("C")
ax1.set_xscale("log")
plt.title("Score vs C for Logistic Regression with StandardScaler")
plt.savefig("Logistic Regression score vs C", bbox_inches = 'tight')
plt.show()

```




```

In [46]: plt.figure()
ax1 = plt.gca()
linearsvc_dataframe = pd.DataFrame(grid_linearsvc.cv_results_)
line1, = ax1.plot(linearsvc_dataframe['param_C'], linearsvc_dataframe['mean_train_score'])
line2, = ax1.plot(linearsvc_dataframe['param_C'], linearsvc_dataframe['mean_test_score'])

plt.fill_between(linearsvc_dataframe.param_C.astype(np.float),
                 linearsvc_dataframe['mean_train_score'] + linearsvc_dataframe['std_train_score'],
                 linearsvc_dataframe['mean_train_score'] - linearsvc_dataframe['std_train_score'],
                 color='orange')
plt.fill_between(linearsvc_dataframe.param_C.astype(np.float),
                 linearsvc_dataframe['mean_test_score'] + linearsvc_dataframe['std_test_score'],
                 linearsvc_dataframe['mean_test_score'] - linearsvc_dataframe['std_test_score'],
                 color='gray')

plt.legend([line1, line2], ["mean_train_score", "mean_test_score"], loc="upper left")
ax1.set_ylabel("score")
ax1.set_xlabel("C")
ax1.set_xscale("log")
plt.title("Score vs C for Linear SVC with StandardScaler")
plt.savefig("Linear SVC score vs C", bbox_inches = 'tight')
plt.show()

```




```

In [52]: plt.figure()
ax1 = plt.gca()
nearest_centroid_dataframe = pd.DataFrame(grid_nearest_centroid.cv_results_)
line1, = ax1.plot(nearest_centroid_dataframe['param_shrink_threshold'], nearest_centroid_dataframe['mean_train_score'])
line2, = ax1.plot(nearest_centroid_dataframe['param_shrink_threshold'], nearest_centroid_dataframe['mean_test_score'])

plt.fill_between(nearest_centroid_dataframe.param_shrink_threshold.astype(np.float64), nearest_centroid_dataframe['mean_train_score'] + nearest_centroid_dataframe['mean_test_score'], nearest_centroid_dataframe['mean_train_score'] - nearest_centroid_dataframe['mean_test_score'])
plt.fill_between(nearest_centroid_dataframe.param_shrink_threshold.astype(np.float64), nearest_centroid_dataframe['mean_test_score'] + nearest_centroid_dataframe['mean_train_score'], nearest_centroid_dataframe['mean_test_score'] - nearest_centroid_dataframe['mean_train_score'])

plt.legend([line1, line2], ["mean_train_score", "mean_test_score"], loc="upper right")
ax1.set_ylabel("score")
ax1.set_xlabel("shrink_threshold")
ax1.set_xscale("log")
plt.title("Score vs shrink_threshold for Nearest Centroid with StandardScaler")
plt.savefig("Nearest Centroid score vs C", bbox_inches = 'tight')
plt.show()

```



## 2.4

```
In [55]: # Kfold for Logistic Regression
from sklearn.model_selection import KFold
params = [{'C': [0.01, 1, 10, 100, 1000]}]
kf = KFold(shuffle=True, random_state = 0)
grid_logreg_kfold = GridSearchCV(LogisticRegression(multi_class='multinomial'),
grid_logreg_kfold.fit(X_train_scaled,y_train)
pd.DataFrame(grid_logreg_kfold.cv_results_)
```

```
Out[55]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	params	rank_test_score
0	13.819211	0.026531	0.722734	0.722888	0.01	{'C': 0.01}	
1	13.279655	0.025883	0.723951	0.724030	1	{'C': 1}	
2	13.156065	0.026231	0.723930	0.724046	10	{'C': 10}	
3	13.596971	0.026013	0.723680	0.723847	100	{'C': 100}	
4	13.549845	0.026424	0.723705	0.723838	1000	{'C': 1000}	

```
In [56]: from sklearn.model_selection import KFold
params = [{'C': [0.01, 1, 10, 100, 1000]}]
kf = KFold(shuffle=True, random_state = 23)
grid_logreg_kfold = GridSearchCV(LogisticRegression(multi_class='multinomial'),
grid_logreg_kfold.fit(X_train_scaled,y_train)
pd.DataFrame(grid_logreg_kfold.cv_results_)
```

```
Out[56]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	params	rank_test_score
0	14.881437	0.029927	0.722764	0.723051	0.01	{'C': 0.01}	
1	13.973566	0.029743	0.723671	0.724040	1	{'C': 1}	
2	13.885630	0.028339	0.723762	0.724038	10	{'C': 10}	
3	15.164577	0.028938	0.723845	0.724104	100	{'C': 100}	
4	14.234138	0.025845	0.723836	0.724089	1000	{'C': 1000}	

```
In [57]: X_train_new, X_test_new, y_train_new, y_test_new = train_test_split(
        forest_data, cover_type, random_state=24)
#scaler = StandardScaler()
scaler.fit(X_train_new)
X_train_new_scaled = scaler.transform(X_train_new)
X_test_new_scaled = scaler.transform(X_test_new)
```



```
In [58]: params = [{'C': [0.01, 1, 10, 100, 1000]}]
kf = KFold(shuffle=True, random_state = 23)
grid_logreg_kfold_new_scaled = GridSearchCV(LogisticRegression(multi_class='
grid_logreg_kfold_new_scaled.fit(X_train_new_scaled,y_train)
pd.DataFrame(grid_logreg_kfold_new_scaled.cv_results_)
```

```
Out[58]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	params	rank_tes
0	12.749484	0.024265	0.487389	0.487438	0.01	{'C': 0.01}	
1	12.939968	0.024612	0.487386	0.487437	1	{'C': 1}	
2	13.010075	0.024633	0.487386	0.487436	10	{'C': 10}	
3	13.607497	0.024616	0.487386	0.487437	100	{'C': 100}	
4	13.669615	0.027509	0.487386	0.487437	1000	{'C': 1000}	

```
In [ ]:
```

```
In [60]: # Kfold for Linear SVC
from sklearn.model_selection import KFold
params = [{'C': [0.01, 1, 10, 100, 1000]}]
kf = KFold(shuffle=True, random_state = 0)
grid_linesvc_kfold = GridSearchCV(LinearSVC(dual=False, tol=0.001), params
grid_linesvc_kfold.fit(X_train_scaled,y_train)
pd.DataFrame(grid_linesvc_kfold.cv_results_)
```

```
Out[60]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	params	rank_tes
0	38.502326	0.036306	0.712297	0.712397	0.01	{'C': 0.01}	
1	43.159197	0.028188	0.712607	0.712786	1	{'C': 1}	
2	44.027135	0.033794	0.712605	0.712793	10	{'C': 10}	
3	45.255886	0.036994	0.712600	0.712787	100	{'C': 100}	
4	42.431367	0.029711	0.712607	0.712793	1000	{'C': 1000}	

In [61]:

```

params = [{'C': [0.01, 1, 10, 100, 1000]}]
kf = KFold(shuffle=True, random_state = 23)
grid_linearsvc_kfold = GridSearchCV(LinearSVC(dual=False, tol=0.001), params)
grid_linearsvc_kfold.fit(X_train_scaled,y_train)
pd.DataFrame(grid_linearsvc_kfold.cv_results_)

```

Out[61]:

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	params	rank_test_score
0	38.096436	0.029992	0.712316	0.712480	0.01	{'C': 0.01}	
1	41.421902	0.029080	0.712605	0.712834	1	{'C': 1}	
2	41.952436	0.037596	0.712616	0.712849	10	{'C': 10}	
3	40.118394	0.031115	0.712614	0.712850	100	{'C': 100}	
4	42.095113	0.027708	0.712612	0.712849	1000	{'C': 1000}	

In [62]:

```

params = [{'C': [0.01, 1, 10, 100, 1000]}]
kf = KFold(shuffle=True, random_state = 23)
grid_linearsvc_kfold_new_scaled = GridSearchCV(LinearSVC(dual=False, tol=0.001), params)
grid_linearsvc_kfold_new_scaled.fit(X_train_new_scaled,y_train)
pd.DataFrame(grid_linearsvc_kfold_new_scaled.cv_results_)

```

Out[62]:

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_C	params	rank_test_score
0	30.087380	0.026045	0.487386	0.487439	0.01	{'C': 0.01}	
1	23.956915	0.026351	0.487386	0.487437	1	{'C': 1}	
2	25.441062	0.032273	0.487386	0.487437	10	{'C': 10}	
3	25.502872	0.033514	0.487386	0.487437	100	{'C': 100}	
4	24.968369	0.028285	0.487386	0.487437	1000	{'C': 1000}	

```
In [65]: # Kfold for Nearest Centroid
params = [{'shrink_threshold': [0,0.5,1,10, 15, 20, 50, 100]}]
kf = KFold(shuffle=True, random_state = 0)
grid_nearest_centroid_kfold = GridSearchCV(NearestCentroid(), params, cv=kf)
grid_nearest_centroid_kfold.fit(X_train_scaled,y_train)
pd.DataFrame(grid_nearest_centroid_kfold.cv_results_)
```

```
Out[65]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_shrink_threshold
0	0.195823	0.038003	0.549905	0.550000	0 {'sr
1	0.395064	0.046341	0.548634	0.548724	0.5 {'sr
2	0.398716	0.035100	0.547112	0.547431	1 {'sr
3	0.377510	0.035181	0.542933	0.543051	10 {'sr
4	0.369522	0.038149	0.547571	0.547768	15 {'sr
5	0.360834	0.041314	0.553226	0.553234	20 {'sr
6	0.354040	0.038796	0.634119	0.634307	50 {'sr
7	0.358631	0.034428	0.515140	0.515244	100 {'sr

```
In [66]: params = [{'shrink_threshold': [0,0.5,1,10, 15, 20, 50, 100]}]
kf = KFold(shuffle=True, random_state = 23)
grid_nearest_centroid_kfold = GridSearchCV(NearestCentroid(), params, cv=kf)
grid_nearest_centroid_kfold.fit(X_train_scaled,y_train)
pd.DataFrame(grid_nearest_centroid_kfold.cv_results_)
```

```
Out[66]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_shrink_threshold
0	0.181291	0.036835	0.549948	0.549852	0 {'sr
1	0.426806	0.039336	0.548735	0.548708	0.5 {'sr
2	0.350931	0.034518	0.547264	0.547296	1 {'sr
3	0.349073	0.033866	0.542974	0.542902	10 {'sr
4	0.338475	0.035174	0.547314	0.547286	15 {'sr
5	0.346626	0.033742	0.552748	0.552847	20 {'sr
6	0.343126	0.033329	0.633919	0.633970	50 {'sr
7	0.346839	0.033320	0.514897	0.515001	100 {'sr

```
In [67]: params = [{'shrink_threshold': [0,0.5,1,10, 15, 20, 50, 100]}]
kf = KFold(shuffle=True, random_state = 23)
grid_nearest_centroid_kfold_new_scaled = GridSearchCV(NearestCentroid(), pa
grid_nearest_centroid_kfold_new_scaled.fit(X_train_new_scaled,y_train)
pd.DataFrame(grid_nearest_centroid_kfold_new_scaled.cv_results_)
```

```
Out[67]:
```

	mean_fit_time	mean_score_time	mean_test_score	mean_train_score	param_shrink_threshold
0	0.202094	0.044167	0.035937	0.037689	0 {'st
1	0.408201	0.038316	0.032440	0.034003	0.5 {'st
2	0.422890	0.044169	0.040481	0.040756	1 {'st
3	0.423583	0.045564	0.364612	0.364613	10 {'st
4	0.383388	0.037520	0.364612	0.364613	15 {'st
5	0.396229	0.034246	0.364612	0.364613	20 {'st
6	0.362487	0.033999	0.364612	0.364613	50 {'st
7	0.369381	0.034582	0.364612	0.364613	100 {'st

## 2.5


```
In [82]: shape = grid_logreg.best_estimator_.coef_.shape
len(grid_logreg.best_estimator_.coef_[0])
```

```
Out[82]: 54
```

For Logistic Regression, stratified Kfold with C=100 and standard scaling gave the best score and hence visualizing the coefficients for that model

```
In [85]: for i in range(0,7):
          plt.scatter(range(54), grid_logreg.best_estimator_.coef_[i],
                      c=np.sign(grid_logreg.best_estimator_.coef_[i]), cmap="b2g")
          plt.title("Coefficients of Logistic Regression with stratified Kfold and C=100 for class 1")
          plt.savefig("Coefficients of Logistic Regression with stratified Kfold and C=100 for class 1")
          plt.show()
```

Coefficients of Logistic Regression with stratified Kfold and C=100 for class 1




Coefficients of Logistic Regression with stratified Kfold and C=100 for class 2



For Linear SVC, stratified Kfold with C=10 and standard scaling gave the best score and hence visualizing the coefficients for that model

```
In [86]: for i in range(0,7):
    plt.scatter(range(54), grid_linearsvc.best_estimator_.coef_[i],
                c=np.sign(grid_linearsvc.best_estimator_.coef_[i]), cmap=
    plt.title("Coefficients of Linear SVC with stratified Kfold and C=10 for
    plt.savefig("Coefficients of Linear SVC with stratified Kfold and C=10 f
    plt.show()
```

Coefficients of Linear SVC with stratified Kfold and C=10 for class 1



Coefficients of Linear SVC with stratified Kfold and C=10 for class 2



```
In [103]: count = 0
j=0
k=0
fig = plt.figure(figsize=(10,50))
for i in range(0,14):
    fig.add_subplot(14,1,i+1)
    if i%2 == 0:
        plt.scatter(range(54), grid_logreg.best_estimator_.coef_[j],
                    c=np.sign(grid_logreg.best_estimator_.coef_[j]), cmap="bwr")
        plt.title("Coefficients of Logistic Regression with stratified Kfold")
        j+=1
    else:
        plt.scatter(range(54), grid_linearsvc.best_estimator_.coef_[k],
                    c=np.sign(grid_linearsvc.best_estimator_.coef_[k]), cmap="bwr")
        plt.title("Coefficients of Linear SVC with stratified Kfold and C=10")
        k+=1

    plt.tight_layout()
    count += 1
plt.show()

<matplotlib.figure.Figure at 0x12db2ca90>
```

Logistic Regression and Linear SVM are agreeing on which features are important

In [ ]: